

Introduction

DS 5110: Big Data Systems

Summer 2026

Lecture 1a

Yue Cheng



UNIVERSITY
of VIRGINIA

Introduction – Yue Cheng

- On the faculty of Data Science & Computer Science
 - Web: <https://tddg.github.io>
 - Email: mrz7dp@virginia.edu
- Current research: Designing **better AI systems**
 - Sustainable AI model storage
 - GPU-efficient Jupyter notebooks
 - Sustainable infra for generative content
 - Faster LLM fine-tuning / serving
 - ...



Course staff and getting help

- Instructor: Yue Cheng
 - Office hours: Thursday 3pm – 4pm, Zoom
- GTAs:
 - Yuyang Cheng
 - Email: jrm9ga@virginia.edu
 - Office hours: Tuesday 2pm – 5pm, Zoom
 - Parvati Viswanathan
 - Email: bsg5sec@virginia.edu
 - Office hours: Wednesday 2pm – 5pm, Zoom



Course staff and getting help


- Offline discussion, questions: Ed
 - <https://edstem.org/us/dashboard>
 - Alternative place to ask questions about assignments, materials, and ideas
 - No anonymous posts or questions
 - Can use private posts to instructor/GTA
 - We are monitoring Ed several times a day
 - We will respond to questions in a batch manner

Today's agenda

- What is this course about?
- What will you do in this course?
- Shell, TUI, and AI coding CLI

A brief history about Big Data


Google Trends



A brief history about Big Data


Google Trends

— MapReduce — Apache Spark




A brief history about Big Data


Google Trends




A brief history about Big Data



A brief history about Big Data



From Big Data to AI boom



Google circa 1997



Search Stanford

10 results ▾

clustering on ▾

Search

Search The Web

10 results ▾

clustering on ▾

Search



Everything is about data

“... **Storage space** must be used efficiently to store indices and, optionally, the documents themselves. The indexing system must process **hundreds of gigabytes** of data efficiently...”


“The system... downloading the last 11 million pages in just 63 hours... The sorter can be **run completely in parallel**; using four machines, the whole process of sorting takes about **24 hours**...”

The anatomy of a large-scale hypertextual Web search engine ¹

Sergey Brin ², Lawrence Page ^{*,2}

Computer Science Department, Stanford University, Stanford, CA 94305, USA

Google circa 2000



Commodity CPUs

Lots of disks

Low bandwidth network

Cheap!






Google

A Google Datacenter in Hamina, Finland



Data explosion in GenAI era



Hugging Face AI model hub

LLM training checkpoints are gigantic!




Checkpoint size = model weights + **optimizer state**

*: <https://mlcommons.org/2025/08/storage-2-checkpointing/>


LLM training checkpoints are gigantic!

A 16,000-GPU cluster has a MTTF of 186 minutes



A trainer should checkpoint 150+ times a day to mitigate the high cost of frequent failures at 16,000-GPU scale

Moore's law is ending



SOURCE: RAY KURZWEIL, "THE SINGULARITY IS NEAR: WHEN HUMANS TRANSCEND BIOLOGY", P.67, THE VIKING PRESS, 2006. DATAPPOINTS BETWEEN 2000 AND 2012 REPRESENT BCA ESTIMATES.

Increased complexity – Computation

Software




CPU

Increased complexity – Computation

Software



CPU



Software




CPU

Increased complexity – Computation

Software



CPU



Software



CPU



GPU



FPGA



ASIC

Increased complexity – More and more choices in clouds

Basic tier: A0, A1, A2, A3, A4
Optimized Compute : D1, D2, D3, D4, D11, D12, D13
D1v2, D2v2, D3v2, D11v2,...
Latest CPUs: G1, G2, G3, ...
Network Optimized: A8, A9
Compute Intensive: A10, A11,...

Microsoft Azure

t2.nano, t2.micro, t2.small
m4.large, m4.xlarge, m4.2xlarge, m4.4xlarge, m3.medium, c4.large, c4.xlarge, c4.2xlarge, c3.large, c3.xlarge, c3.4xlarge, r3.large, r3.xlarge, r3.4xlarge, i2.2xlarge, i2.4xlarge, d2.xlarge, d2.2xlarge, d2.4xlarge,...

Amazon EC2

n1-standard-1, ns1-standard-2, ns1-standard-4, ns1-standard-8, ns1-standard-16, ns1-highmem-2, ns1-highmem-4, ns1-highmem-8, n1-highcpu-2, n1-highcpu-4, n1-highcpu-8, n1-highcpu-16, n1-highcpu-32, f1-micro, g1-small...

Google Cloud



But how do we program this to tackle the challenges of big data in the AI era?



Course syllabus

Big picture course goals

- Learn about some of the most influential works in (big) data + AI systems
- Explain the design and architecture
- Read and evaluate some seminal papers
- Develop and deploy applications on open-source data systems and public cloud (AWS, Google Cloud)
- Design and report some data/AI systems ideas

Schedule (tentative)

- Slides, assignments, due dates
- Less concrete further out; don't get too far ahead

<https://tddg.github.io/ds5110-summer26/>

The screenshot shows the course page for DS5110: Big Data Systems, Summer 2026. The page has a dark blue header with navigation links: Syllabus, Schedule, Materials, Assignments, and Staff. The main content area is dark blue with white text. It features the course title 'DS5110: Big Data Systems' in large font, followed by a subtitle 'Scalable data systems, parallel analytics, cloud infrastructure, and modern AI systems/infrastructure'. Below this are three boxes containing course details: Meeting Time (Mon-Fri, 9AM to 11:15AM), Instructor (Yue Cheng), and Location (Clark Hall 107). At the bottom, there is a yellow box for 'LATEST ANNOUNCEMENT' dated May 18, 2026, with the text 'Welcome' and 'Welcome to DS 5110!'.

DS5110

Syllabus Schedule Materials Assignments Staff

SUMMER 2026

DS5110: Big Data Systems

Scalable data systems, parallel analytics, cloud infrastructure, and modern AI systems/infrastructure

MEETING TIME
Mon-Fri, 9AM to 11:15AM

INSTRUCTOR
Yue Cheng

LOCATION
Clark Hall 107

LATEST ANNOUNCEMENT

Announcements

MAY 18, 2026
Welcome
Welcome to DS 5110!

Lectures (tentative schedule)

- Lecture (+ demos + labs)
 - Slides available on course website (night before)
- Week 1: Basics, MapReduce, Spark + AI tools
- Week 2: Project ideas + cloud data systems
- Week 3: AI systems / infrastructure
- Week 4: Advanced research topics in AI systems + final project presentations

Textbooks?

- Papers, documentations, blog articles (required or optional) serve as reference for many topics that aren't directly covered by a text
- Slides/lecture notes
- Three optional textbooks (first two are free)
 - **“Operating Systems: Three Easy Pieces (OSTEP)”** by Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau
 - **“Distributed Systems (3rd edition)”** by van Steen and Tenenbaum will supply optional alternate explanations
 - **“Designing Data-Intensive Applications (1st edition)”** by Martin Kleppmann (can be accessed via UVA library)

- 2.5 hands-on assignments, with **AI coding CLI (opencode)**, on **AWS**
 - **Assignment 0:** AWS Academy, EC2, Linux shell, and AI coding CLI
 - **Assignment 1:** Analytics with Apache Spark and DuckDB
 - **Assignment 2:** Burst-parallel ML with AWS Lambda
- Short coding assignments
 - Gain hands-on experience with AWS, popular open-source data systems and cloud tools
 - Get exposed to AI coding CLIs
- Assignments and project are all team-based

Grading


- Assignments (50% total)
 - Assignment 0 (10%)
 - Assignment 1 (20%)
 - Assignment 2 (20%)
- Project (30%)
- Final presentation (20%)

Background survey

Background survey

What courses have you taken before


58 responses



Background survey

Which programming languages or query languages are you most comfortable with?


58 responses



Background survey

Which AI coding/development tools have you used?


58 responses



Background survey

How often do you currently use AI coding tools?


58 responses



Background survey

For which tasks do you primarily use AI tools?

58 responses



Time to introduce yourself