

AI Storage Infrastructure

DS5110: Big Data Systems

Summer 2026

Lecture 10

Yue Cheng



UNIVERSITY
of VIRGINIA

Material taken/derived from:

- Harvard ML Systems by Vijay Janapa Reddi

@ 2026 released for use under a [CC BY-SA](#) license.

Workload inversion

AI inverts every storage assumption

Traditional (Database)

Access Pattern Small random I/O (4 KB reads)
Working Set Hot data fits in cache (80/20 rule)
Key Metric IOPS (operations per second)
RAID Strategy RAID 5/6 (redundancy over bandwidth)
Write Pattern Continuous small writes (WAL)

ML Training

Access Pattern Large sequential streaming (MB reads)
Working Set Entire dataset read every epoch
Key Metric Throughput (GB/s sustained)
RAID Strategy RAID 0 (bandwidth over redundancy)
Write Pattern Massive checkpoint bursts

The I/O wall: When storage throughput cannot deliver training data as fast as accelerators consume it, GPUs idle regardless of their computational power.

The small file problem

Storing one file per sample is catastrophic at scale:

- Each file: `open()`, `stat()`, `close()`
- Metadata server: 100K—300K ops/s
- 10,000 workers → saturates in <1 sec

	Small Files	Shards
Files	14M	1,400
Metadata ops	42M/epoch	4,200/epoch
Read pattern	Random	Sequential
Throughput	0.5 GB/s	25 GB/s

The small file problem

Storing one file per sample is catastrophic at scale:


- Each file: `open()`, `stat()`, `close()`
- Metadata server: 100K—300K ops/s
- 10,000 workers → saturates in <1 sec

	Small Files	Shards
Files	14M	1,400
Metadata ops	42M/epoch	4,200/epoch
Read pattern	Random	Sequential
Throughput	0.5 GB/s	25 GB/s

Solution: Aggregate into large sequential shards (TFRecords, WebDataset, Parquet).
Reduce metadata ops by 10,000—100,000x

Storage hierarchy

The six-tier AI storage hierarchy



Each tier trades bandwidth for capacity and cost. Physics of distance governs the tradeoff.

Design principle: Right data, right tier, right time. Each tier is a prefetch buffer for the tier above.

Bandwidth cliffs between tiers

Tier	Storage	Bandwidth	Cliff
0	GPU HBM	3.35 TB/s	↓ 17×
1	Host DRAM	200 GB/s	↓ 15×
2	Local NVMe	7–25 GB/s	↓ 3–10×
3	Parallel FS	4 GB/s/node	↓ 10×
4	Object Storage	~1 GB/s/node	

Implication: HBM empties in 24 ms but takes 400 ms to refill from DRAM. The next batch must *already* be in HBM before the current one finishes. Every tier prefetches for the tier above.

Tier 0: GPU HBM – The destination

HBM is the only tier where weights and activations reside during computation.

Component (per GPU)	Size
Weights (FP16, tensor-parallel)	43.75 GB
Optimizer (ZeRO-3)	5.5 GB
Activations	10–20 GB
Gradient buffers	5.5 GB
NCCL buffers	2–4 GB
Total occupied	67–79 GB

H100: 80GB HBM

85—98% occupied

<15% free for prefetch

Late batch = GPU stall

Early batch = no room

Just-in-time, not just-too-late

Predict: Text vs. image bandwidth

Q: A 2,048-GPU cluster with 200 ms iteration time

Estimate the aggregate storage bandwidth needed for:

- a) Text** training (4,096 tokens/GPU, 4 bytes/token)
- b) Image** training (256 images/GPU, 150 KB/image)

Reveal: Text vs. image bandwidth

	Text	Image	Ratio
Batch size/GPU	4,096 tokens $\times$ 4 B	256 imgs $\times$ 150 KB	
Per-GPU data	16 KB	38.4 MB	2,400 $\times$
Aggregate BW	160 MB/s	393 GB/s	2,400 $\times$
Storage needed	Single NAS node	High-perf PFS	

Text: bandwidth-light, volume-heavy. Bottleneck = dataset size & checkpoint

Image: bandwidth-heavy. Bottleneck = storage system's ability to feed accelerators

Pipeline equation

The data pipeline equation

$$BW_{\text{req}} = N_{GPUs} \times U_{\text{target}} \times \frac{S_{\text{batch}}}{T_{\text{iteration}}}$$

Pipeline equation: Four levers

Variable	Meaning
N_{GPUs}	Number of accelerators
U_{target}	Target utilization (0.8–0.95)
S_{batch}	Local batch size (bytes)
$T_{iteration}$	Forward-backward time

Key: No lever is free. Reducing any input has a cost.

Worked example (ImageNet, 256 GPUs):

$$\begin{aligned} & 256 \times 0.80 \times \frac{256 \times 150KB}{0.2s} \\ &= 256 \times 0.80 \times 192 \text{ MB/s} \\ &= \mathbf{39.3 \text{ GB/s}} \end{aligned}$$

Must sustain **39.3 GB/s** continuously for days.
Fall short → accelerators idle.

The data stall ratio

$$\text{Data Stall \%} = \frac{T_{step} - T_{compute}}{T_{step}} \times 100$$


Stall Ratio	Diagnosis	Action
< 2%	Healthy pipeline	Monitor, do not over-invest
2–10%	Marginal (will break at scale)	Audit pipeline, add prefetch
> 10%	Clear storage bottleneck	Redesign data path

Economic consequence: 5% stall on 1,000 GPUs at \$2/GPU-hr for 1 week = **\$16,800 wasted**. Storage monitoring is a first-order economic concern.


Pipelining and prefetching

Pipelined Data Loading: Hiding Storage Latency

Without Pipelining (Sequential)




With Pipelining (Overlapped)



Pipelining: 450 ms to 250 ms (1.8x faster)

Pipelining and prefetching

Pipelined Data Loading: Hiding Storage Latency



Critical caveat: Pipelining hides average latency. Buffer depth must absorb **P99 tail latency** to prevent stalls **at scale**.

Quick check

What is the 500x storage—compute gap, and what technique hides it?

Size the prefetch buffer

Size the prefetch buffer


A training node has:

- $T_{\text{compute}} = 200$ ms per batch
- $T_{\text{I/O,avg}} = 250$ ms (avg load time)
- $T_{\text{I/O, P99}} = 500$ ms (tail load latency)

Calculate min prefetch depth (in batches) for
(a) average and (b) P99 latency.

Economics

The storage cost iceberg



Counterintuitive: The cheapest storage tier is not the cheapest choice. Total cost of data delivery includes egress + idle GPUs.

Worked example: Stream vs. Cache

50 TB vision dataset, 20 epochs, 4 training runs/year

Option A: Stream from S3

Storage: \$13,800/yr

Egress: $50 \text{ TB} \times 20 \times \$0.09/\text{GB} = \$90\text{K}/\text{run}$

4 runs/yr: \$360K egress

Total: \$373,800/yr

Option B: NVMe cache

S3 storage: \$13,800/yr

NVMe capacity: \$60,000/yr

Egress (stage 1x/run): \$18,000/yr

Total: \$91,800/yr

NVMe caching saves \$282,000/year: enough for 3 additional H100 GPUs. Break-even: 4 epochs of reading per training run. Multi-epoch training always benefits from local caching.

Tiering strategy


Tier	Storage	What Lives Here	\$/GB/mo
Hot	Local NVMe + DRAM	Active training data, current batch	\$0.10
Warm	Parallel FS	Shared datasets, recent checkpoints	\$0.03
Cold	Object (S3)	All organizational data, versioned	\$0.02
Archive	Glacier	Compliance, reproducibility	\$0.004

Lifecycle: Data promotes upward when a job requests it, demotes downward when idle. Automate with policies, not engineers.

Staging patterns: *Pre-stage* (best steady-state, slow start), *On-demand* (slow first epoch), *Streaming* (zero startup, no cache). Choice depends on epoch count and local NVMe capacity.

AI failures and checkpoints

AI failure spectrum



Fault tolerance must address all three categories: each requires different detection and recovery mechanisms.

Key: Hard failures are easy. Silent data corruption is the real threat at scale.

Hardware fault taxonomy

Fault Type	Behavior	Rate	Detection
Transient	Vanishes on retry (bit flip)	1 per 10K GPU-h	Silent
Permanent	Component dead (GPU, NIC, disk)	1 per 50K GPU-h	ECC/heartbeat
Intermittent	Works 99% of the time	1 per 25K GPU-h	Very hard
Byzantine	Sends incorrect data while “healthy”	1 per 100K GPU-h	Requires voting

At 10K GPUs: ~24 transient faults/day, 5 permanent/day. Every category occurs daily at scale.

Transient fault mechanism

Space

Galactic cosmic rays

Solar cosmic rays

Atmosphere

① Reactions in the atmosphere

High-energy particle

High-energy particle breaks nucleus

Nucleus (Nitrogen, Oxygen)

proton

neutron

π -meson

Muon

② Reactions in telecommunication equipment

Soft error

Temporary error (Recover by restart/overwrite etc.)


Electrical noises

Semiconductor device
in telecommunication
equipment

Neutron breaks
silicon nucleus


Silicon nucleus

Creating
secondary particles



Faults and the bathtub curve

The Bathtub Curve: Failure Rate Over System Lifetime



Most ML training runs experience infant mortality failures -- they rarely run long enough for wear-out

Economic case: \$10M/month cluster $\times$ 30% downtime = \$3M/month recoverable via fault tolerance.

The mathematics of inevitable failure

Individual component reliability:

$$R_{single}(t) = e^{-\lambda t}$$

For N identical components (series):

$$R_{system}(t) = e^{-N\lambda t}$$

The mathematics of inevitable failure

Individual component reliability:

$$R_{single}(t) = e^{-\lambda t}$$

For N identical components (series):

$$R_{system}(t) = e^{-N\lambda t}$$


System MTBF:

$$MTBF_{system} = \frac{MTBF_{component}}{N}$$

Scale inverts reliability:

GPUs	Cluster MTBF
8	260 days
1,000	2 days
10,000	5 hours
25,000	2 hours
Per-GPU MTBF: 50,000 hours	

The reliability gap



Assumes per-GPU MTBF = 100,000 hours (server-grade). System MTBF = MTBF_component / N.

Implication: At 10K+ GPUs, fault tolerance is not optional – it is **physically necessary**.

Quick check

At 10,000 GPUs with 50,000h per-GPU MTBF, what is the cluster MTBF?

Predict: The 9s of reliability

Each GPU has **99.999%** uptime.


You have **10,000 GPUs**.

What is the probability the entire cluster survives **one hour** without a failure?

The Young-Daly Law

$$\tau_{\text{opt}} = \sqrt{2 \cdot T_{\text{write}} \cdot \text{MTBF}}$$

Optimal checkpoint interval



Young-Daly: Worked example

Scenario: Llama-3 training on 16,000 H100 GPUs

Parameter	Value	Source
Checkpoint write time (T_w)	2 min	Storage BW
Cluster MTBF	3 hrs (180 min)	16K GPUs

$$\tau_{\text{opt}} = \sqrt{2 \cdot 2 \cdot 180} = \sqrt{720} \approx \mathbf{27 \text{ min}}$$

Checkpoint every 10 min: 22.7% overhead (too frequent)


Checkpoint every 60 min: 20% overhead (too risky)

Checkpoint every 27 min: 14.9% overhead (optimal)

Checkpoint storage


Checkpoint Staging: Minimizing T_{save}

Training Timeline with Checkpoint Storm



All 256 nodes write
1,050 GB simultaneously

Tiered Checkpoint Staging (Decouples T_{save} from PFS)



175B Model Checkpoint Math

Weights (FP16): 350 GB
Optimizer state (FP32): 1,400 GB
Total per checkpoint: 1,050 GB (with ZeRO)
30-day run: ~4.5 PB of checkpoint writes

Staging Impact

Without staging: $T_{\text{save}} = 150$ sec (PFS limited)
With NVMe staging: $T_{\text{save}} = 15$ sec
Training pause reduced by 10x
Async replication: zero impact on training

Checkpoint storm: All nodes write simultaneously. PFS must be provisioned for peak burst, not average. Write locally, replicate asynchronously.

The compute data path

Data lifecycle during a 30-day training run (checkpoint every 10 min):

Phase	Data Path	Volume
Setup	Object → NVMe (stage once)	3 TB
Training	NVMe → DRAM → HBM (per batch)	3 TB/epoch
Checkpoint	HBM → NVMe (local, fast)	1,050 GB each
Replicate	NVMe → PFS (async, background)	1,050 GB each
Archive	PFS → Object → Glacier (lifecycle)	~50 TB retained

Counterintuitive: Checkpoint I/O (~4.5 PB) dwarfs training data I/O (3 TB per epoch). For frontier models, **checkpoint staging strategy** has more impact on storage design than training data pipeline optimization.